

HW 1 - MLP and Character Language Modeling-4

January 30, 2023

1 Homework 1 - MLP and Character Language Modeling

1.1 Name: Yichen Huang

1.2 CUID: yh3550

```
[ ]: import torch
from torch.utils.data import DataLoader
from torch.utils.data.dataset import random_split
import torch.nn as nn
from torch.utils.data import Dataset
from torch.utils.data import DataLoader, TensorDataset
import time
from tqdm import tqdm
```

1.2.1 Information

- We will do a few preliminary exercises and also build a character level MLP language model.
- This model will be similar to the model we did in class, except that we will have characters as tokens, not words.
- You will need a conda environment for this, here is general information on this.
- <https://docs.conda.io/projects/conda/en/latest/user-guide/install/index.html>
- PyTorch: <https://anaconda.org/pytorch/pytorch>

In the code below, FILL-IN the code necessary in the hint string provided.

1.2.2 Preliminary exercises

- Please fill in the cells below with the asked for data.

```
[ ]: torch.manual_seed(1)
```

```
[ ]: <torch._C.Generator at 0x108433bf0>
```

```
[ ]: # Create an embedding layer for a vocabulary of size 10 and the word vectors_
      ↪are each of dimension 5.
e = nn.Embedding(10, 5)
```

```

# Extract the embedding for the word whose token index is 3. What is the shape
↳ of this vector?
v = e(torch.tensor(3))
print("Shape of the vector: ", v.size())

# Extract the weight matrix from the layer e.
# Create a linear layer (with no bias) of size 10 by 5 and set it's data to the
↳ embedding matrix.
l = nn.Linear(5, 10, bias=False)
l.weight = e.weight

# Insert inside of the assert below some sort of equality check between l.
↳ weight and e.weight; it should pass to true.
# Hint: look up torch.all() and torch.eq()

assert(torch.eq(l.weight, e.weight).all())

```

Shape of the vector: torch.Size([5])

```

[ ]: # Create a batch of size 2 with entries [0, 1, 2] and [2, 3, 4] in the data
↳ batch.
x = torch.tensor([[0, 1, 2], [2, 3, 4]])

```

```

[ ]: # What is the dimension of this batch ran through the embedding layer?
assert(e(x).shape == torch.Size([2, 3, 5]))

```

```

[ ]:

```

1.2.3 Constants and configs used below.

```

[ ]: DEVICE = "cpu"
LR = 4.0
BATCH_SIZE = 16
NUM_EPOCHS = 5
MARKER = '.'
# N-gram level;  $P(w_t \mid w_{t-1}, \dots, w_{t-n+1})$ .
# We use 3 words to predict the next word.
n = 4
# Hidden layer dimension.
h = 20
# Word embedding dimension.
m = 20

```

```

[ ]:

```

1.2.4 Get the dataset and the tokenizer.

```
[ ]: class CharDataset(Dataset):
    def __init__(self, words, chars):
        self.words = words
        self.chars = chars
        # Inverse dictionaries mapping char tokens to unique ids and the
        ↪reverse.
        # Tokens in this case are the unique chars we passed in above.
        # Each token should be mappend to a unique integer and MARKER should
        ↪have token 0.
        # For example, stoi should be like {'.' -> 0, 'a' -> 1, 'b' -> 2} if I
        ↪pass in chars = '.ab'.
        dic_stoi = {}
        dic_itos = {}

        for i in range(len(chars)):
            dic_stoi[chars[i]] = i
            dic_itos[i] = chars[i]

        self.stoi = dic_stoi
        self.itos = dic_itos

    def __len__(self):
        # Number of words.
        return len(self.words)

    def contains(self, word):
        # Check if word is in self.words and return True/False if it is, is not.
        return word in self.words

    def get_vocab_size(self):
        # Return the vocabulary size.
        return len(self.chars)

    def encode(self, word):
        # Express this word as a list of int ids. For example, maybe ".abc" ->
        ↪[0, 1, 2, 3].
        # This assumes 'a' -> 1, etc.
        return [self.stoi[char] for char in word]

    def decode(self, tokens):
        # For a set of tokens, return back the string.
        # For example, maybe [1, 1, 2] -> "aac"
        return ''.join([self.itos[token] for token in tokens])

    def __getitem__(self, idx):
```

```

    # This is used so we can loop over the data.
    word = self.words[idx]
    return self.encode(word)

```

```
[ ]:
```

```

[ ]: def create_datasets(window, input_file = 'names.txt'):
    """
    This takes a file of words and separates all the words.
    It then gets all the characters present in the universe of words and then
    outputs the statistics.
    """
    with open(input_file, 'r') as f:
        data = f.read()
    # Split the file by new lines. You should get a list of names.
    words = data.split('\n')

    words = [word.replace(' ', '') for word in words] # This gets rid of any
    trailing and starting white spaces.
    words = [i for i in words if i] # Filter out all the empty words.

    # This gets the universe of all characters.
    chars = sorted(list(set([char for word in words for char in word])))

    # Will force chars to have MARKER having index 0.
    chars = [MARKER] + chars

    # Pad each word with a context window of size n-1.
    # Why? a word like "abc" should becomes "..abc.." if the window is size 3.
    # This is some we can get pair of (x, y) data like this: ".." -> "a", ".a"
    -> "b", "ab" -> "c", "bc" -> ".", "c." -> "."
    # I.e. this allows us to know that "a" is a start character.
    # So you should get something like ["ab", "c"] -> ["..ab..", "..c.."], for
    example.
    words = [MARKER * (window - 1) + word + MARKER * (window - 1) for word in
    words]

    print(f"The number of examples in the dataset: {len(words)}")
    print(f"The number of unique characters in the vocabulary: {len(chars)}")
    print(f"The vocabulary we have is: {''.join(chars)}")

    # Partition the input data into a training, validation, and the test set.
    out_of_sample_set_size = min(2000, int(len(words) * 0.1)) # We use 10% of
    the training set, or up to 2000 examples.
    test_set_size = 1500

```

```

# First, get a random permutation of randomly permute of size len(words).
# Then, convert this to a list.
# This index list is used below to get the train, validation, and test sets.
rp = torch.randperm(len(words)).tolist()

# Get train, validation, and test set.
train_words = [words[i] for i in rp[:-out_of_sample_set_size]]
validation_words = [words[i] for i in rp[-out_of_sample_set_size:
↪-test_set_size]]
test_words = [words[i] for i in rp[-test_set_size:]]

print(f"We've split up the dataset into {len(train_words)}, ↵
↪{len(validation_words)}, {len(test_words)} training, validation, and test ↵
↪examples")

# But the data in the data set objects.
train_dataset = CharDataset(train_words, chars)
validation_dataset = CharDataset(validation_words, chars)
test_dataset = CharDataset(test_words, chars)

return train_dataset, validation_dataset, test_dataset

```

```
[ ]: train_dataset, validation_dataset, test_dataset = create_datasets(n)
```

The number of examples in the dataset: 32033

The number of unique characters in the vocabulary: 27

The vocabulary we have is: .abcdefghijklmnopqrstuvwxyz

We've split up the dataset into 30033, 500, 1500 training, validation, and test examples

```
[ ]: train_dataset.words[0]
```

```
[ ]: '...niyam...'
```

1.3 Explore the data

```
[ ]: # Get the first word in "train_dataset"
train_dataset.words[0]
```

```
[ ]: '...niyam...'
```

```
[ ]:
```

```
[ ]: # Get the stoi map of train_dataset. How many keys does it have?
len(train_dataset.stoi)
```

```
[ ]: 27
```

```
[ ]:
```

1.3.1 Get the dataloader

```
[ ]: def create_dataloader(dataset, window):
    x_list = []
    y_list = []
    # For each word.
    for i, word in enumerate(dataset):
        # Grab a context of size window and window-1 characters will be in x, 1
        # character will be in y.
        for j, _ in enumerate(word):
            # If there is no window of size window left, break.
            if j + window > len(word) - 1:
                break
            word_window = word[j:j+window]
            x, y = word_window[:window-1], word_window[-1]
            x_list.append(x)
            y_list.append(y)

    return DataLoader(
        TensorDataset(torch.tensor(x_list), torch.tensor(y_list)),
        BATCH_SIZE,
        shuffle=True
    )
```

```
[ ]:
```

```
[ ]: train_dataloader = create_dataloader(train_dataset, n)
validation_dataloader = create_dataloader(validation_dataset, n)
test_dataloader = create_dataloader(test_dataset, n)
```

```
[ ]:
```

1.3.2 Set up the model

- Identical to lecture. Please look over that!

```
[ ]: # One of the first Neural language models!
class CharacterNeuralLanguageModel(nn.Module):
    def __init__(self, V, m, h, n):
        super(CharacterNeuralLanguageModel, self).__init__()

        # Vocabulary size.
        self.V = V

        # Embedding dimension, per word.
```

```

self.m = m

# Hidden dimension.
self.h = h

# N in "N-gram"
self.n = n

# Can you change all this stuff to use nn.Linear?
# Ca also use nn.Parameter(torch.zeros(V, m)) for self.C but then we
↪ need one-hot and this is slow.
self.C = nn.Embedding(V, m)
self.H = nn.Parameter(torch.zeros((n-1) * m, h))
self.W = nn.Parameter(torch.zeros((n-1) * m, V))
self.U = nn.Parameter(torch.zeros(h, V))

self.b = torch.nn.Parameter(torch.ones(V))
self.d = torch.nn.Parameter(torch.ones(h))

self.init_weights()

def init_weights(self):
    # Initialize C, H, W, U in a nice way. Use xavier initialization for
    ↪ the weights.
    # On a first run, just pass.
    with torch.no_grad():
        torch.nn.init.xavier_uniform_(self.C.weight)
        torch.nn.init.xavier_uniform_(self.H)
        torch.nn.init.xavier_uniform_(self.W)
        torch.nn.init.xavier_uniform_(self.U)

def forward(self, x):

    # x is of dimension N = batch size X n-1

    # N X (n-1) X m
    x = self.C(x)

    # N
    N = x.shape[0]

    # N X (n-1) * m
    x = x.view(N, -1)

    # N X V

```

```

        y = self.b + torch.matmul(x, self.W) + torch.matmul(nn.Tanh()(self.d +
↪torch.matmul(x, self.H)), self.U)

    return y

```

```
[ ]:
```

1.3.3 Set up the model.

```

[ ]: # Identical to lecture.
criterion = torch.nn.CrossEntropyLoss().to(DEVICE)
model = CharacterNeuralLanguageModel(
    train_dataset.get_vocab_size(), m, h, n
).to(DEVICE)
optimizer = torch.optim.SGD(model.parameters(), lr=LR)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, 1.0, gamma=0.1)

[ ]: # How many parameters does the neural network have?
# Hint: look up model.named_parameters and the method "nelement" on a tensor.
# See also the XOR notebook where we count the gradients that are 0.
# There, we loop over the parameters.
number_parameters = 0
for name, param in model.named_parameters():
    print(f"Parameter {name} has {param.numel()} elements.")
    number_parameters += param.numel()
print(f"Total number of parameters: {number_parameters}")

```

```

Parameter H has 1200 elements.
Parameter W has 1620 elements.
Parameter U has 540 elements.
Parameter b has 27 elements.
Parameter d has 20 elements.
Parameter C.weight has 540 elements.
Total number of parameters: 3947

```

1.3.4 Train the model.

```

[ ]: def calculate_perplexity(total_loss, total_batches):
    return torch.exp(torch.tensor(total_loss / total_batches)).item()

[ ]: def train(dataloader, model, optimizer, criterion, epoch):
    model.train()
    total_loss, total_batches = 0.0, 0.0
    log_interval = 500

    for idx, (x, y) in tqdm(enumerate(dataloader)):
        optimizer.zero_grad()

```



```

logits = model(x)

# Get the loss.
loss = criterion(input=logits, target=y.view(-1))

# Do back propagation.
loss.backward()

# Clip the gradients so they don't explode. Look at how this is done in
↳ lecture.
torch.nn.utils.clip_grad_norm_(model.parameters(), 0.1)

# Do an optimization step.
optimizer.step()

total_loss += loss.item()
total_batches += 1

if idx % log_interval == 0 and idx > 0:
    perplexity = calculate_perplexity(total_loss, total_batches)
    print(
        "| epoch {:3d} "
        "| {:5d}/{:5d} batches "
        "| perplexity {:8.3f} "
        "| loss {:8.3f} "
        .format(
            epoch,
            idx,
            len(dataloader),
            perplexity,
            total_loss / total_batches,
        )
    )
    total_loss, total_batches = 0.0, 0

```

```

[ ]: def evaluate(dataloader, model, criterion):
    model.eval()
    total_loss, total_batches = 0.0, 0

    with torch.no_grad():
        for idx, (x, y) in enumerate(dataloader):
            logits = model(x)
            total_loss += criterion(input=logits, target=y.squeeze(-1)).item()
            total_batches += 1
    return total_loss / total_batches, calculate_perplexity(total_loss,
↳ total_batches)

```

```

[ ]: for epoch in range(1, NUM_EPOCHS + 1):
    epoch_start_time = time.time()
    train(train_dataloader, model, optimizer, criterion, epoch)
    loss_val, perplexity_val = evaluate(validation_dataloader, model, criterion)
    scheduler.step()
    print("-" * 59)
    print(
        "| end of epoch {:3d} |"
        "| time: {:5.2f}s |"
        "| valid perplexity {:8.3f} |"
        "| valid loss {:8.3f}".format(
            epoch,
            time.time() - epoch_start_time,
            perplexity_val,
            loss_val
        )
    )
    print("-" * 59)

print("Checking the results of test dataset.")
loss_test, perplexity_test = evaluate(test_dataloader, model, criterion)
print("test perplexity {:8.3f} | test loss {:8.3f} ".format(perplexity_test,
↵loss_test))

```

867it [00:00, 1879.58it/s]

epoch	1		500/15247 batches		perplexity	9.705		loss	2.273
-------	---	--	-------------------	--	------------	-------	--	------	-------

1332it [00:00, 2120.97it/s]

epoch	1		1000/15247 batches		perplexity	8.738		loss	2.168
-------	---	--	--------------------	--	------------	-------	--	------	-------

1762it [00:00, 2002.05it/s]

epoch	1		1500/15247 batches		perplexity	8.605		loss	2.152
-------	---	--	--------------------	--	------------	-------	--	------	-------

2450it [00:01, 2201.15it/s]

epoch	1		2000/15247 batches		perplexity	8.323		loss	2.119
-------	---	--	--------------------	--	------------	-------	--	------	-------

2927it [00:01, 2290.79it/s]

epoch	1		2500/15247 batches		perplexity	8.510		loss	2.141
-------	---	--	--------------------	--	------------	-------	--	------	-------

3421it [00:01, 2380.11it/s]

epoch	1		3000/15247 batches		perplexity	8.326		loss	2.119
-------	---	--	--------------------	--	------------	-------	--	------	-------

epoch	1		3500/15247 batches		perplexity	8.231		loss	2.108
-------	---	--	--------------------	--	------------	-------	--	------	-------

4388it [00:02, 2209.57it/s]

epoch	1		4000/15247 batches		perplexity	8.015		loss	2.081
-------	---	--	--------------------	--	------------	-------	--	------	-------

4850it [00:02, 2258.03it/s]

epoch	1	4500/15247 batches	perplexity	8.167	loss	2.100
5311it [00:02, 2251.25it/s]						
epoch	1	5000/15247 batches	perplexity	8.002	loss	2.080
5780it [00:02, 2281.83it/s]						
epoch	1	5500/15247 batches	perplexity	8.121	loss	2.094
6230it [00:02, 2181.59it/s]						
epoch	1	6000/15247 batches	perplexity	8.067	loss	2.088
6911it [00:03, 2232.63it/s]						
epoch	1	6500/15247 batches	perplexity	8.083	loss	2.090
7362it [00:03, 2106.52it/s]						
epoch	1	7000/15247 batches	perplexity	7.990	loss	2.078
7800it [00:03, 2070.10it/s]						
epoch	1	7500/15247 batches	perplexity	7.733	loss	2.045
8283it [00:03, 2243.23it/s]						
epoch	1	8000/15247 batches	perplexity	8.021	loss	2.082
8732it [00:04, 2204.82it/s]						
epoch	1	8500/15247 batches	perplexity	8.104	loss	2.092
9389it [00:04, 2107.83it/s]						
epoch	1	9000/15247 batches	perplexity	7.951	loss	2.073
9814it [00:04, 2069.37it/s]						
epoch	1	9500/15247 batches	perplexity	7.866	loss	2.063
10231it [00:04, 2062.58it/s]						
epoch	1	10000/15247 batches	perplexity	7.756	loss	2.049
10633it [00:05, 1737.98it/s]						
epoch	1	10500/15247 batches	perplexity	8.197	loss	2.104
11273it [00:05, 2001.55it/s]						
epoch	1	11000/15247 batches	perplexity	8.008	loss	2.080
11926it [00:05, 2126.75it/s]						
epoch	1	11500/15247 batches	perplexity	7.859	loss	2.062
12433it [00:05, 2341.58it/s]						
epoch	1	12000/15247 batches	perplexity	7.530	loss	2.019
epoch	1	12500/15247 batches	perplexity	7.986	loss	2.078

```

13345it [00:06, 2142.49it/s]
| epoch   1 | 13000/15247 batches | perplexity   8.037 | loss   2.084
13781it [00:06, 2152.49it/s]
| epoch   1 | 13500/15247 batches | perplexity   7.898 | loss   2.067
14213it [00:06, 2096.74it/s]
| epoch   1 | 14000/15247 batches | perplexity   7.566 | loss   2.024
14848it [00:07, 2015.59it/s]
| epoch   1 | 14500/15247 batches | perplexity   7.897 | loss   2.066
15247it [00:07, 2118.47it/s]
| epoch   1 | 15000/15247 batches | perplexity   7.815 | loss   2.056
-----
| end of epoch   1 | time:  7.23s | valid perplexity   7.777 | valid loss
2.051
-----

763it [00:00, 2557.36it/s]
| epoch   2 |  500/15247 batches | perplexity   7.514 | loss   2.017
| epoch   2 | 1000/15247 batches | perplexity   7.111 | loss   1.962
1812it [00:00, 2557.69it/s]
| epoch   2 | 1500/15247 batches | perplexity   7.164 | loss   1.969
| epoch   2 | 2000/15247 batches | perplexity   7.318 | loss   1.990
2868it [00:01, 2623.91it/s]
| epoch   2 | 2500/15247 batches | perplexity   7.256 | loss   1.982
| epoch   2 | 3000/15247 batches | perplexity   7.209 | loss   1.975
3940it [00:01, 2533.90it/s]
| epoch   2 | 3500/15247 batches | perplexity   7.180 | loss   1.971
4482it [00:01, 2622.21it/s]
| epoch   2 | 4000/15247 batches | perplexity   7.157 | loss   1.968
| epoch   2 | 4500/15247 batches | perplexity   7.158 | loss   1.968
5307it [00:02, 2707.60it/s]
| epoch   2 | 5000/15247 batches | perplexity   7.254 | loss   1.982
| epoch   2 | 5500/15247 batches | perplexity   7.012 | loss   1.948
6400it [00:02, 2723.80it/s]
| epoch   2 | 6000/15247 batches | perplexity   6.993 | loss   1.945
| epoch   2 | 6500/15247 batches | perplexity   7.224 | loss   1.977
7502it [00:02, 2734.41it/s]

```

epoch	2	7000/15247 batches	perplexity	7.133	loss	1.965
epoch	2	7500/15247 batches	perplexity	7.176	loss	1.971

8324it [00:03, 2723.46it/s]

epoch	2	8000/15247 batches	perplexity	7.039	loss	1.951
epoch	2	8500/15247 batches	perplexity	7.034	loss	1.951

9435it [00:03, 2765.71it/s]

epoch	2	9000/15247 batches	perplexity	7.378	loss	1.998
epoch	2	9500/15247 batches	perplexity	6.957	loss	1.940

10545it [00:03, 2765.15it/s]

epoch	2	10000/15247 batches	perplexity	7.155	loss	1.968
epoch	2	10500/15247 batches	perplexity	7.195	loss	1.973

11378it [00:04, 2767.66it/s]

epoch	2	11000/15247 batches	perplexity	7.045	loss	1.952
epoch	2	11500/15247 batches	perplexity	7.052	loss	1.953

12491it [00:04, 2753.19it/s]

epoch	2	12000/15247 batches	perplexity	6.973	loss	1.942
epoch	2	12500/15247 batches	perplexity	7.191	loss	1.973

13322it [00:04, 2754.72it/s]

epoch	2	13000/15247 batches	perplexity	7.003	loss	1.946
epoch	2	13500/15247 batches	perplexity	7.090	loss	1.959

14434it [00:05, 2733.21it/s]

epoch	2	14000/15247 batches	perplexity	7.184	loss	1.972
epoch	2	14500/15247 batches	perplexity	7.057	loss	1.954

15247it [00:05, 2696.19it/s]

epoch	2	15000/15247 batches	perplexity	7.010	loss	1.947
-------	---	---------------------	------------	-------	------	-------

end of epoch	2	time: 5.68s	valid perplexity	7.023	valid loss	1.949
--------------	---	-------------	------------------	-------	------------	-------

814it [00:00, 2723.84it/s]

epoch	3	500/15247 batches	perplexity	6.910	loss	1.933
epoch	3	1000/15247 batches	perplexity	6.954	loss	1.939

1924it [00:00, 2758.81it/s]

epoch	3	1500/15247 batches	perplexity	7.033	loss	1.951
epoch	3	2000/15247 batches	perplexity	7.119	loss	1.963

3021it [00:01, 2691.37it/s]

epoch	3	2500/15247 batches	perplexity	7.005	loss	1.947
epoch	3	3000/15247 batches	perplexity	6.869	loss	1.927
3851it [00:01, 2739.73it/s]						
epoch	3	3500/15247 batches	perplexity	7.055	loss	1.954
epoch	3	4000/15247 batches	perplexity	6.900	loss	1.932
4957it [00:01, 2719.72it/s]						
epoch	3	4500/15247 batches	perplexity	6.883	loss	1.929
epoch	3	5000/15247 batches	perplexity	7.012	loss	1.948
5781it [00:02, 2726.22it/s]						
epoch	3	5500/15247 batches	perplexity	7.241	loss	1.980
epoch	3	6000/15247 batches	perplexity	6.970	loss	1.942
6882it [00:02, 2717.79it/s]						
epoch	3	6500/15247 batches	perplexity	7.060	loss	1.954
epoch	3	7000/15247 batches	perplexity	7.143	loss	1.966
7983it [00:02, 2729.84it/s]						
epoch	3	7500/15247 batches	perplexity	6.882	loss	1.929
epoch	3	8000/15247 batches	perplexity	6.963	loss	1.941
8814it [00:03, 2753.37it/s]						
epoch	3	8500/15247 batches	perplexity	6.906	loss	1.932
epoch	3	9000/15247 batches	perplexity	7.049	loss	1.953
9928it [00:03, 2777.36it/s]						
epoch	3	9500/15247 batches	perplexity	7.069	loss	1.956
epoch	3	10000/15247 batches	perplexity	7.055	loss	1.954
11041it [00:04, 2777.08it/s]						
epoch	3	10500/15247 batches	perplexity	6.960	loss	1.940
epoch	3	11000/15247 batches	perplexity	6.884	loss	1.929
11879it [00:04, 2784.07it/s]						
epoch	3	11500/15247 batches	perplexity	6.933	loss	1.936
epoch	3	12000/15247 batches	perplexity	7.222	loss	1.977
12998it [00:04, 2783.73it/s]						
epoch	3	12500/15247 batches	perplexity	6.977	loss	1.943
epoch	3	13000/15247 batches	perplexity	7.151	loss	1.967
13837it [00:05, 2785.63it/s]						
epoch	3	13500/15247 batches	perplexity	7.082	loss	1.958
epoch	3	14000/15247 batches	perplexity	7.020	loss	1.949
14956it [00:05, 2786.33it/s]						

epoch	3	14500/15247 batches	perplexity	7.114	loss	1.962
epoch	3	15000/15247 batches	perplexity	7.068	loss	1.956

15247it [00:05, 2737.15it/s]

end of epoch	3	time: 5.60s	valid perplexity	6.994	valid loss	1.945
--------------	---	-------------	------------------	-------	------------	-------

268it [00:00, 2675.88it/s]

epoch	4	500/15247 batches	perplexity	6.875	loss	1.928
-------	---	-------------------	------------	-------	------	-------

825it [00:00, 2759.31it/s]

epoch	4	1000/15247 batches	perplexity	7.027	loss	1.950
-------	---	--------------------	------------	-------	------	-------

1379it [00:00, 2761.86it/s]

epoch	4	1500/15247 batches	perplexity	7.130	loss	1.964
-------	---	--------------------	------------	-------	------	-------

1935it [00:00, 2761.45it/s]

epoch	4	2000/15247 batches	perplexity	7.178	loss	1.971
-------	---	--------------------	------------	-------	------	-------

2491it [00:00, 2762.72it/s]

epoch	4	2500/15247 batches	perplexity	7.017	loss	1.948
-------	---	--------------------	------------	-------	------	-------

3048it [00:01, 2772.53it/s]

epoch	4	3000/15247 batches	perplexity	6.986	loss	1.944
-------	---	--------------------	------------	-------	------	-------

3327it [00:01, 2775.84it/s]

epoch	4	3500/15247 batches	perplexity	7.015	loss	1.948
-------	---	--------------------	------------	-------	------	-------

3884it [00:01, 2771.06it/s]

epoch	4	4000/15247 batches	perplexity	7.055	loss	1.954
-------	---	--------------------	------------	-------	------	-------

4444it [00:01, 2777.06it/s]

epoch	4	4500/15247 batches	perplexity	6.887	loss	1.930
-------	---	--------------------	------------	-------	------	-------

4999it [00:01, 2767.72it/s]

epoch	4	5000/15247 batches	perplexity	6.832	loss	1.922
-------	---	--------------------	------------	-------	------	-------

5557it [00:02, 2774.43it/s]

epoch	4	5500/15247 batches	perplexity	7.057	loss	1.954
-------	---	--------------------	------------	-------	------	-------

5836it [00:02, 2777.75it/s]

epoch	4	6000/15247 batches	perplexity	7.094	loss	1.959
-------	---	--------------------	------------	-------	------	-------

6393it [00:02, 2770.89it/s]

epoch	4	6500/15247 batches	perplexity	7.051	loss	1.953
-------	---	--------------------	------------	-------	------	-------

```

6949it [00:02, 2755.90it/s]
| epoch    4 | 7000/15247 batches | perplexity    7.015 | loss    1.948
7502it [00:02, 2756.45it/s]
| epoch    4 | 7500/15247 batches | perplexity    6.942 | loss    1.938
7780it [00:02, 2761.22it/s]
| epoch    4 | 8000/15247 batches | perplexity    6.986 | loss    1.944
8333it [00:03, 2730.65it/s]
| epoch    4 | 8500/15247 batches | perplexity    6.943 | loss    1.938
8887it [00:03, 2748.87it/s]
| epoch    4 | 9000/15247 batches | perplexity    6.885 | loss    1.929
9439it [00:03, 2749.48it/s]
| epoch    4 | 9500/15247 batches | perplexity    6.844 | loss    1.923
9992it [00:03, 2750.26it/s]
| epoch    4 | 10000/15247 batches | perplexity    7.073 | loss    1.956
10545it [00:03, 2736.83it/s]
| epoch    4 | 10500/15247 batches | perplexity    6.989 | loss    1.944
10819it [00:03, 2732.23it/s]
| epoch    4 | 11000/15247 batches | perplexity    7.023 | loss    1.949
11373it [00:04, 2752.56it/s]
| epoch    4 | 11500/15247 batches | perplexity    6.851 | loss    1.924
11926it [00:04, 2749.79it/s]
| epoch    4 | 12000/15247 batches | perplexity    6.905 | loss    1.932
12480it [00:04, 2755.61it/s]
| epoch    4 | 12500/15247 batches | perplexity    6.861 | loss    1.926
13036it [00:04, 2762.33it/s]
| epoch    4 | 13000/15247 batches | perplexity    6.967 | loss    1.941
13313it [00:04, 2755.84it/s]
| epoch    4 | 13500/15247 batches | perplexity    7.020 | loss    1.949
13867it [00:05, 2757.71it/s]
| epoch    4 | 14000/15247 batches | perplexity    7.028 | loss    1.950
14422it [00:05, 2761.70it/s]
| epoch    4 | 14500/15247 batches | perplexity    6.986 | loss    1.944

```



```

14981it [00:05, 2768.03it/s]
| epoch   4 | 15000/15247 batches | perplexity   7.137 | loss    1.965
15247it [00:05, 2753.05it/s]
-----
| end of epoch   4 | time:  5.57s | valid perplexity   6.987 | valid loss
1.944
-----

541it [00:00, 2710.51it/s]
| epoch   5 |   500/15247 batches | perplexity   7.017 | loss    1.948
817it [00:00, 2728.82it/s]
| epoch   5 |  1000/15247 batches | perplexity   7.090 | loss    1.959
1362it [00:00, 2674.27it/s]
| epoch   5 |  1500/15247 batches | perplexity   7.062 | loss    1.955
1906it [00:00, 2701.16it/s]
| epoch   5 |  2000/15247 batches | perplexity   6.987 | loss    1.944
2461it [00:00, 2740.81it/s]
| epoch   5 |  2500/15247 batches | perplexity   7.005 | loss    1.947
3013it [00:01, 2742.31it/s]
| epoch   5 |  3000/15247 batches | perplexity   7.160 | loss    1.969
3288it [00:01, 2733.97it/s]
| epoch   5 |  3500/15247 batches | perplexity   7.224 | loss    1.977
3841it [00:01, 2748.02it/s]
| epoch   5 |  4000/15247 batches | perplexity   6.921 | loss    1.935
4395it [00:01, 2758.84it/s]
| epoch   5 |  4500/15247 batches | perplexity   7.065 | loss    1.955
4947it [00:01, 2742.78it/s]
| epoch   5 |  5000/15247 batches | perplexity   6.747 | loss    1.909
5497it [00:02, 2717.40it/s]
| epoch   5 |  5500/15247 batches | perplexity   7.056 | loss    1.954
6047it [00:02, 2730.58it/s]
| epoch   5 |  6000/15247 batches | perplexity   6.975 | loss    1.942
6321it [00:02, 2729.65it/s]
| epoch   5 |  6500/15247 batches | perplexity   6.820 | loss    1.920

```

6880it [00:02, 2760.48it/s]
| epoch 5 | 7000/15247 batches | perplexity 6.858 | loss 1.925
7434it [00:02, 2758.06it/s]
| epoch 5 | 7500/15247 batches | perplexity 7.035 | loss 1.951
7988it [00:02, 2758.74it/s]
| epoch 5 | 8000/15247 batches | perplexity 6.927 | loss 1.935
8542it [00:03, 2762.34it/s]
| epoch 5 | 8500/15247 batches | perplexity 6.937 | loss 1.937
8819it [00:03, 2755.54it/s]
| epoch 5 | 9000/15247 batches | perplexity 6.960 | loss 1.940
9371it [00:03, 2743.91it/s]
| epoch 5 | 9500/15247 batches | perplexity 7.090 | loss 1.959
9922it [00:03, 2738.61it/s]
| epoch 5 | 10000/15247 batches | perplexity 6.949 | loss 1.939
10480it [00:03, 2762.04it/s]
| epoch 5 | 10500/15247 batches | perplexity 6.912 | loss 1.933
11037it [00:04, 2764.92it/s]
| epoch 5 | 11000/15247 batches | perplexity 6.918 | loss 1.934
11314it [00:04, 2763.43it/s]
| epoch 5 | 11500/15247 batches | perplexity 6.971 | loss 1.942
11867it [00:04, 2746.71it/s]
| epoch 5 | 12000/15247 batches | perplexity 6.879 | loss 1.928
12423it [00:04, 2761.01it/s]
| epoch 5 | 12500/15247 batches | perplexity 6.909 | loss 1.933
12980it [00:04, 2764.75it/s]
| epoch 5 | 13000/15247 batches | perplexity 7.007 | loss 1.947
13534it [00:04, 2758.92it/s]
| epoch 5 | 13500/15247 batches | perplexity 7.042 | loss 1.952
13810it [00:05, 2757.07it/s]
| epoch 5 | 14000/15247 batches | perplexity 7.005 | loss 1.947
14363it [00:05, 2757.52it/s]
| epoch 5 | 14500/15247 batches | perplexity 7.007 | loss 1.947

14920it [00:05, 2768.93it/s]

| epoch 5 | 15000/15247 batches | perplexity 6.977 | loss 1.943

15247it [00:05, 2742.18it/s]

| end of epoch 5 | time: 5.59s | valid perplexity 6.974 | valid loss 1.942

Checking the results of test dataset.

test perplexity 7.121 | test loss 1.963

Hint: For the above, you should see your loss around 2.0 and going down. Similarly to perplexity which should be around 7 to 8.

1.4 Generate some text.

```
[ ]: def generate_word(model, dataset, window):
    generated_word = []
    # Set the context to a window-1 length array having just the MARKER
    ↪ character's token_id.
    context = (window - 1) * [dataset.stoi[MARKER]]

    while True:
        logits = model(torch.tensor(context).view(1, -1))

        # Get the probabilities from the logits.
        # Hint: softmax!
        probs = nn.Softmax(dim=1)(logits)

        # Get 1 sample from a multinomial having the above probabilities.
        token_id = torch.multinomial(probs, 1).item()

        # Append the token_id to the generated word.
        generated_word.append(token_id)

        # Move the context over 1, drop the first (oldest) token and append the
    ↪ new one above.
        # The size of the resulting context should be the same.
        # For example, if it was "[0, 1, 2]" and you generated 4, it should now
    ↪ be [1, 2, 4].
        context = context[1:] + [token_id]

        if token_id == 0:
            # If you generate token_id = 0, i.e. '.', break out.
            break

    # Return and decode the generated word to a string.
    return dataset.decode(generated_word)
```

```
[ ]: torch.manual_seed(1)
      for _ in range(50):
          print(generate_word(model, train_dataset, n))
```

ama.
ele.
lia.
aldi.
jarorsse.
dez.
bria.
jairestlei.
revy.
madlais.
hoanna.
dacelian.
alalie.
shais.
maya.
jouston.
zailah.
ede.
rie.
gros.
aukh.
bamaka.
anyaarius.
kelee.
har.
jami.
naekshreem.
kaylen.
quyla.
naygusen.
mayanatram.
ahazoriexsunya.
shamonti.
hori.
evfiah.
rosie.
vaivel.
ynalaydin.
kenasia.
dar.
wun.
jayana.
ris.

nor.
ilyn.
marri.
alevante.
kalyn.
desleeshanaa.
daniellaenimariinilah.